

# Storage

## Living Standard — Last Updated 26 February 2021



### Participate:

[GitHub whatwg/storage](#) ([new issue](#), [open issues](#))  
[IRC: #whatwg on Freenode](#)

### Commits:

[GitHub whatwg/storage/commits](#)  
[Snapshot as of this commit](#)  
[@storagestandard](#)

### Tests:

[web-platform-tests storage/](#) ([ongoing work](#))

### Translations (non-normative):

[日本語](#)

# Abstract

The Storage Standard defines an API for persistent storage and quota estimates, as well as the platform storage architecture.

## Table of Contents

<a href="#">1 Introduction</a>	
<a href="#">2 Terminology</a>	
<a href="#">3 Lay of the land</a>	
<a href="#">4 Model</a>	
<a href="#">4.1 Storage endpoints</a>	
<a href="#">4.2 Storage keys</a>	
<a href="#">4.3 Storage sheds</a>	
<a href="#">4.4 Storage shelves</a>	
<a href="#">4.5 Storage buckets</a>	
<a href="#">4.6 Storage bottles</a>	
<a href="#">4.7 Storage proxy maps</a>	
<a href="#">5 Persistence permission</a>	
<a href="#">6 Usage and quota</a>	
<a href="#">7 Management</a>	
<a href="#">7.1 Storage pressure</a>	
<a href="#">7.2 User interface guidelines</a>	
<a href="#">8 API</a>	
<a href="#">Acknowledgments</a>	
<a href="#">Intellectual property rights</a>	
<a href="#">Index</a>	
<a href="#">Terms defined by this specification</a>	
<a href="#">Terms defined by reference</a>	
<a href="#">References</a>	
<a href="#">Normative References</a>	
<a href="#">IDL Index</a>	

## § 1. Introduction

Over the years the web has grown various APIs that can be used for storage, e.g., IndexedDB, localStorage, and showNotification(). The Storage Standard consolidates these APIs by defining:

- A bucket, the primitive these APIs store their data in
- A way of making that bucket persistent
- A way of getting usage and quota estimates for an [origin](#)

Traditionally, as the user runs out of storage space on their device, the data stored with these APIs gets lost without the user being able to intervene. However, persistent buckets cannot be cleared without consent by the user. This thus brings data guarantees users have enjoyed on native platforms to the web.

### Example

A simple way to make storage persistent is through invoking the [persist\(\)](#) method. It simultaneously requests the end user for permission and changes the storage to be persistent once granted:

```
navigator.storage.persist().then(persisted => {
  if (persisted) {
    /* ... */
  }
});
```

To not show user-agent-driven dialogs to the end user unannounced slightly more involved code can be written:

```
Promise.all([
  navigator.storage.persist(),
  navigator.permissions.query({name: "persistent-storage"})
]).then(([persisted, permission]) => {
  if (!persisted && permission.state == "granted") {
    navigator.storage.persist().then( /* ... */ );
  } else if (!persisted && permission.state == "prompt") {
    showPersistentStorageExplanation();
  }
});
```

The [estimate\(\)](#) method can be used to determine whether there is enough space left to store content for an application:

```
function retrieveNextChunk(nextChunkInfo) {
  return navigator.storage.estimate().then(info => {
    if (info.quota - info.usage > nextChunkInfo.size) {
      return fetch(nextChunkInfo.url);
    } else {
      throw new Error("insufficient space to store next chunk");
    }
  }).then( /* ... */ );
}
```

## § 2. Terminology

This specification depends on the Infra Standard. [\[INFRA\]](#)

This specification uses terminology from the HTML, IDL, and Permissions Standards.  
[\[HTML\]](#) [\[WEBIDL\]](#) [\[PERMISSIONS\]](#)

## § 3. Lay of the land

A [user agent](#) has various kinds of semi-persistent state:

### **Credentials**

End-user credentials, such as username and passwords submitted through HTML forms

### **Permissions**

Permissions for various features, such as geolocation

### **Network**

HTTP cache, cookies, authentication entries, TLS client certificates

### **Storage**

Indexed DB, Cache API, service worker registrations, `localStorage`, `sessionStorage`, application caches, notifications, etc.

This standard primarily concerns itself with storage.

## § 4. Model

Standards defining local or session storage APIs will define a [storage endpoint](#) and [register](#) it by changing this standard. They will invoke either the [obtain a local storage bottle map](#) or the [obtain a session storage bottle map](#) algorithm, which will give them:

- Failure, which might mean the API has to throw or otherwise indicate there is no storage available for that [environment settings object](#).
- A [storage proxy map](#) that operates analogously to a [map](#), which can be used to store data in a manner that suits the API. This standard takes care of isolating that data from other APIs, [storage keys](#), and [storage types](#).

### Note

*If you are defining a standard for such an API, consider filing an issue against this standard for assistance and review.*

To isolate this data this standard defines a [storage shed](#) which segments [storage shelves](#) by a [storage key](#). A [storage shelf](#) in turn consists of a [storage bucket](#) and will likely consist of multiple [storage buckets](#) in the future to allow for different storage policies. And lastly, a [storage bucket](#) consists of [storage bottles](#), one for each [storage endpoint](#).

## § 4.1. Storage endpoints

A **storage endpoint** is a [local](#) or [session storage](#) API that uses the infrastructure defined by this standard, most notably [storage bottles](#), to keep track of its storage needs.

A [storage endpoint](#) has an **identifier**, which is a [storage identifier](#).

A [storage endpoint](#) also has **types**, which is a [set](#) of [storage types](#).

A [storage endpoint](#) also has a **quota**, which is null or a number representing a recommended [quota](#) (in bytes) for each [storage bottle](#) corresponding to this [storage endpoint](#).

A **storage identifier** is an [ASCII string](#).

A **storage type** is "local" or "session".

The **registered storage endpoints** are a [set](#) of [storage endpoints](#) defined by the following table:

<a href="#">Identifier</a>	<a href="#">Type</a>	<a href="#">Quota</a>
"caches"	« "local" »	null
"indexedDB"	« "local" »	null
"localStorage"	« "local" »	$5 \times 2^{20}$ (i.e., 5 mebibytes)
"serviceWorkerRegistrations"	« "local" »	null
"sessionStorage"	« "session" »	$5 \times 2^{20}$ (i.e., 5 mebibytes)

## NOTE

As mentioned, standards can use these [storage identifiers](#) with [obtain a local storage bottle map](#) and [obtain a session storage bottle map](#). It is anticipated that some APIs will be applicable to both [storage types](#) going forward.

## § 4.2. Storage keys

A **storage key** is an [origin](#). [\[HTML\]](#)

This is expected to change; see [Client-Side Storage Partitioning](#).

To **obtain a storage key**, given an [environment settings object](#) *environment*, run these steps:

1. Let *key* be *environment*'s [origin](#).
2. If *key* is an [opaque origin](#), then return failure.
3. If the user has disabled storage, then return failure.
4. Return *key*.

## § 4.3. Storage sheds

A **storage shed** is a [map](#) of [storage keys](#) to [storage shelves](#). It is initially empty.

A [user agent](#) holds a **storage shed**, which is a [storage shed](#). A user agent's [storage shed](#) holds all **local storage** data.

A [browsing session](#) holds a **storage shed**, which is a [storage shed](#). A [browsing session](#)'s [storage shed](#) holds all **session storage** data.

To **legacy-clone a browsing session storage shed**, given a [browsing session](#) *A* and a [browsing session](#) *B*, run these steps:

1. [For each](#) *key* → *shelf* of *A*'s [storage shed](#):
  1. Let *newShelf* be the result of running [create a storage shelf](#) with "session".
  2. Set *newShelf*'s [bucket map](#)["default"]'s [bottle map](#)["sessionStorage"]'s [map](#) to a [clone](#) of *shelf*'s [bucket map](#)["default"]'s [bottle map](#)["sessionStorage"]'s [map](#).
  3. Set *B*'s [storage shed](#)[*key*] to *newShelf*.

## Note

*This is considered legacy as the benefits, if any, do not outweigh the implementation complexity. And therefore it will not be expanded or used outside*

## § 4.4. Storage shelves

A **storage shelf** exists for each [storage key](#) within a [storage shed](#). It holds a **bucket map**, which is a [map](#) of [strings](#) to [storage buckets](#).

Note

*For now "default" is the only [key](#) that exists in a [bucket map](#). See [issue #2](#). It is given a [value](#) when a [storage shelf](#) is [obtained](#) for the first time.*

To **obtain a storage shelf**, given a [storage shed](#) *shed*, an [environment settings object](#) *environment*, and a [storage type](#) *type*, run these steps:

1. Let *key* be the result of running [obtain a storage key](#) with *environment*.
2. If *key* is failure, then return failure.
3. If *shed*[*key*] does not [exist](#), then set *shed*[*key*] to the result of running [create a storage shelf](#) with *type*.
4. Return *shed*[*key*].

To **obtain a local storage shelf**, given an [environment settings object](#) *environment*, return the result of running [obtain a storage shelf](#) with the user agent's [storage shed](#), *environment*, and "local".

To **create a storage shelf**, given a [storage type](#) *type*, run these steps:

1. Let *shelf* be a new [storage shelf](#).
2. Set *shelf*'s [bucket map](#)["default"] to the result of running [create a storage bucket](#) with *type*.
3. Return *shelf*.

## § 4.5. Storage buckets

A **storage bucket** is a place for [storage endpoints](#) to store data.

A [storage bucket](#) has a **bottle map** of [storage identifiers](#) to [storage bottles](#).

A **local storage bucket** is a [storage bucket](#) for [local storage](#) APIs.

A [local storage bucket](#) has a **mode**, which is "best-effort" or "persistent". It is initially "best-effort".

A **session storage bucket** is a [storage bucket](#) for [session storage](#) APIs.



To **create a storage bucket**, given a [storage type](#) type, run these steps:

1. Let *bucket* be null.
2. If *type* is "local", then set *bucket* to a new [local storage bucket](#).
3. Otherwise:
  1. Assert: *type* is "session".
  2. Set *bucket* to a new [session storage bucket](#).
4. **For each** *endpoint* of [registered storage endpoints](#) whose [types contain](#) *type*, set *bucket*'s [bottle map](#)[*endpoint*'s [identifier](#)] to a new [storage bottle](#) whose [quota](#) is *endpoint*'s [quota](#).
5. Return *bucket*.

## § 4.6. Storage bottles

A **storage bottle** is a part of a [storage bucket](#) carved out for a single [storage endpoint](#). A [storage bottle](#) has a **map**, which is initially an empty [map](#). A [storage bottle](#) also has a **proxy map reference set**, which is initially an empty [set](#). A [storage bottle](#) also has a **quota**, which is null or a number representing a conservative estimate of the total amount of bytes it can hold. Null indicates the lack of a limit.

Note *It is still bound by the [storage quota](#) of its encompassing [storage shelf](#).*

A [storage bottle](#)'s [map](#) is where the actual data meant to be stored lives. User agents are expected to store this data, and make it available across [agent](#) and even [agent cluster](#) boundaries, in an [implementation-defined](#) manner, so that this standard and standards using this standard can access the contents.

To **obtain a storage bottle map**, given a [storage type](#) type, [environment settings object](#) *environment*, and [storage identifier](#) *identifier*, run these steps:

1. Let *shed* be null.
2. If *type* is "local", then set *shed* to the user agent's [storage shed](#).
3. Otherwise:
  1. Assert: *type* is "session".
  2. Set *shed* to *environment*'s [browsing session](#)'s [storage shed](#).
4. Let *shelf* be the result of running [obtain a storage shelf](#), with *shed*, *environment*, and *type*.
5. If *shelf* is failure, then return failure.
6. Let *bucket* be *shelf*'s [bucket map](#)["default"].
7. Let *bottle* be *bucket*'s [bottle map](#)[*identifier*].
8. Let *proxyMap* be a new [storage proxy map](#) whose [backing map](#) is *bottle*'s [map](#).
9. [Append](#) *proxyMap* to *bottle*'s [proxy map reference set](#).

10. Return *proxyMap*.

To **obtain a local storage bottle map**, given an [environment settings object](#) *environment* and [storage identifier](#) *identifier*, return the result of running [obtain a storage bottle map](#) with "local", *environment*, and *identifier*.

To **obtain a session storage bottle map**, given an [environment settings object](#) *environment* and [storage identifier](#) *identifier*, return the result of running [obtain a storage bottle map](#) with "session", *environment*, and *identifier*.

## § 4.7. Storage proxy maps

A **storage proxy map** is equivalent to a [map](#), except that all operations are instead performed on its **backing map**.

This allows for the [backing map](#) to be replaced. This is needed for [issue #4](#) and potentially the [Storage Access API](#).

## § 5. Persistence permission

A [local storage bucket](#) can only have its [mode](#) change to "persistent" if the user (or user agent on behalf of the user) has granted permission to use the "[persistent-storage](#)" feature.

### Note

*When granted to an [origin](#), the persistence permission can be used to protect storage from the user agent's clearing policies. The user agent cannot clear storage marked as persistent without involvement from the [origin](#) or user. This makes it particularly useful for resources the user needs to have available while offline or resources the user creates locally.*

The "[persistent-storage](#)" [powerful feature](#)'s permission-related flags, algorithms, and types are defaulted, except for:

### [permission state](#)

"[persistent-storage](#)"'s [permission state](#) must have the same value for all [environment settings objects](#) with a given [origin](#).

### [permission revocation algorithm](#)

1. If "[persistent-storage](#)"'s [permission state](#) is "[granted](#)", then return.
2. Let *shelf* be the result of running [obtain a local storage shelf](#) with [current settings object](#).
3. Set *shelf*'s [bucket map](#)["default"]'s [mode](#) to "best-effort".

## § 6. Usage and quota

The **storage usage** of a [storage shelf](#) is an [implementation-defined](#) rough estimate of the amount of bytes used by it.

### Note

*This cannot be an exact amount as user agents might, and are encouraged to, use deduplication, compression, and other techniques that obscure exactly how much bytes a [storage shelf](#) uses.*

The **storage quota** of a [storage shelf](#) is an [implementation-defined](#) conservative estimate of the total amount of bytes it can hold. This amount should be less than the total storage space on the device. It must not be a function of the available storage space on the device.

### Note

*User agents are strongly encouraged to consider navigation frequency, recency of visits, bookmarking, and [permission](#) for "[persistent-storage](#)" when determining quotas.*

*Directly or indirectly revealing available storage space can lead to fingerprinting and leaking information outside the scope of the [origin](#) involved.*

## § 7. Management

Whenever a [storage bucket](#) is cleared by the user agent, it must be cleared in its entirety. User agents should avoid clearing [storage buckets](#) while script that is able to access them is running, unless instructed otherwise by the user.

If removal of [storage buckets](#) leaves the encompassing [storage shelf](#)'s [bucket map empty](#), then [remove](#) that [storage shelf](#) and corresponding [storage key](#) from the encompassing [storage shed](#).

### § 7.1. Storage pressure

A user agent that comes under storage pressure should clear network state and [local storage buckets](#) whose [mode](#) is "best-effort", ideally prioritizing removal in a manner that least impacts the user.

If a user agent continues to be under storage pressure, then the user agent should inform the user and offer a way to clear the remaining [local storage buckets](#), i.e., those whose [mode](#) is "persistent".

[Session storage buckets](#) must be cleared as [browsing sessions](#) are closed.

Note

*If the user agent allows for revival of [browsing sessions](#), e.g., through reopening [browsing sessions](#) or continued use of them after restarting the user agent, then clearing necessarily involves a more complex set of heuristics.*

### § 7.2. User interface guidelines

User agents should offer users the ability to clear network state and storage for individual websites. User agents should not distinguish between network state and storage in their user interface. This ensures network state cannot be used to revive storage and reduces the number of concepts users need to be mindful of.

Credentials should be separated as they contain data the user might not be able to revive, such as an autogenerated password. Permissions are best separated too to avoid inconveniencing the user.

## § 8. API

```
[SecureContext]
interface mixin NavigatorStorage {
  [SameObject] readonly attribute StorageManager storage;
};
Navigator includes NavigatorStorage;
WorkerNavigator includes NavigatorStorage;
```

Each [environment settings object](#) has an associated [StorageManager](#) object. [\[HTML\]](#)

The **storage** getter steps are to return [this](#)'s [relevant settings object](#)'s [StorageManager](#) object.

```
[SecureContext,
  Exposed=(Window,Worker)]
interface StorageManager {
  Promise<boolean> persisted();
  [Exposed=Window] Promise<boolean> persist();

  Promise<StorageEstimate> estimate();
};

dictionary StorageEstimate {
  unsigned long long usage;
  unsigned long long quota;
};
```

The **persisted()** method steps are:

1. Let *promise* be a new promise.
2. Let *shelf* be the result of running [obtain a local storage shelf](#) with [this](#)'s [relevant settings object](#).
3. If *shelf* is failure, then reject *promise* with a [TypeError](#).
4. Otherwise, run these steps [in parallel](#):
  1. Let *persisted* be true if *shelf*'s [bucket map](#)["default"]'s [mode](#) is "persistent"; otherwise false.

Note

*It will be false when there's an internal error.*

2. [Queue a task](#) to resolve *promise* with *persisted*.

5. Return *promise*.

The **persist()** method steps are:

1. Let *promise* be a new promise.
2. Let *shelf* be the result of running [obtain a local storage shelf](#) with [this](#)'s [relevant settings object](#).
3. If *shelf* is failure, then reject *promise* with a [TypeError](#).

4. Otherwise, run these steps [in parallel](#):

1. Let *permission* be the result of [requesting permission to use "persistent-storage"](#).

Note

*User agents are encouraged to not let the user answer this question twice for the same [origin](#) around the same time and this algorithm is not equipped to handle such a scenario.*

2. Let *bucket* be *shelf*'s [bucket map](#)["default"].

3. Let *persisted* be true if *bucket*'s [mode](#) is "persistent"; otherwise false.

Note

*It will be false when there's an internal error.*

4. If *persisted* is false and *permission* is ["granted"](#), then:

1. Set *bucket*'s [mode](#) to "persistent".
2. If there was no internal error, then set *persisted* to true.

5. [Queue a task](#) to resolve *promise* with *persisted*.

5. Return *promise*.

The [estimate\(\)](#) method steps are:

1. Let *promise* be a new promise.
2. Let *shelf* be the result of running [obtain a local storage shelf](#) with [this](#)'s [relevant settings object](#).
3. If *shelf* is failure, then reject *promise* with a [TypeError](#).
4. Otherwise, run these steps [in parallel](#):
  1. Let *usage* be [storage usage](#) for *shelf*.
  2. Let *quota* be [storage quota](#) for *shelf*.
  3. Let *dictionary* be a new [StorageEstimate](#) dictionary whose [usage](#) member is *usage* and [quota](#) member is *quota*.
  4. If there was an internal error while obtaining *usage* and *quota*, then [queue a task](#) to reject *promise* with a [TypeError](#).

Note

*Internal errors are supposed to be extremely rare and indicate some kind of low-level platform or hardware fault. However, at the scale of the web with the diversity of implementation and platforms, the unexpected does occur.*

5. Otherwise, [queue a task](#) to resolve *promise* with *dictionary*.

5. Return *promise*.

## § Acknowledgments

With that, many thanks to Adrian Bateman, Aislinn Grigas, Alex Russell, Ali Alabbas, Andrew Sutherland, Ben Kelly, Ben Turner, Dale Harvey, David Grogan, Domenic Denicola, fantasai, Jake Archibald, Jeffrey Yasskin, Jesse Mykolyn, Jinho Bang, Jonas Sicking, Joshua Bell, Kenji Baheux, Kinuko Yasuda, Luke Wagner, Michael Nordman, Mounir Lamouri, Shachar Zohar, 黃強 (Shawn Huang), 簡冠庭 (Timothy Guan-tin Chien), and Victor Costan for being awesome!

This standard is written by [Anne van Kesteren](#) ([Mozilla](#), [annevk@annevk.nl](mailto:annevk@annevk.nl)).



## § Intellectual property rights

Copyright © WHATWG (Apple, Google, Mozilla, Microsoft). This work is licensed under a [Creative Commons Attribution 4.0 International License](#).

This is the Living Standard. Those interested in the patent-review version should view the [Living Standard Review Draft](#).

## § Index

### § Terms defined by this specification

- [backing map](#), in §4.7
- [bottle map](#), in §4.5
- [bucket map](#), in §4.4
- [create a storage bucket](#), in §4.5
- [create a storage shelf](#), in §4.4
- [estimate\(\)](#), in §8
- [identifier](#), in §4.1
- [legacy-clone a browsing session storage shed](#), in §4.3
- [local storage](#), in §4.3
- [local storage bucket](#), in §4.5
- [map](#), in §4.6
- [mode](#), in §4.5
- [NavigatorStorage](#), in §8
- [obtain a local storage bottle map](#), in §4.6
- [obtain a local storage shelf](#), in §4.4
- [obtain a session storage bottle map](#), in §4.6
- [obtain a storage bottle map](#), in §4.6
- [obtain a storage key](#), in §4.2
- [obtain a storage shelf](#), in §4.4
- [persist\(\)](#), in §8
- [persisted\(\)](#), in §8
- [proxy map reference set](#), in §4.6
- quota
  - [dfn for storage bottle](#), in §4.6
  - [dfn for storage endpoint](#), in §4.1
  - [dict-member for StorageEstimate](#), in §8
- [registered storage endpoints](#), in §4.1
- [session storage](#), in §4.3
- [session storage bucket](#), in §4.5
- [storage](#), in §8
- [storage bottle](#), in §4.6
- [storage bucket](#), in §4.5
- [storage endpoint](#), in §4.1
- [StorageEstimate](#), in §8
- [storage identifier](#), in §4.1
- [storage key](#), in §4.2
- [StorageManager](#), in §8
- [storage proxy map](#), in §4.7
- [storage quota](#), in §6
- storage shed
  - [definition of](#), in §4.3
  - [dfn for browsing session](#), in §4.3
  - [dfn for user agent](#), in §4.3
- [storage shelf](#), in §4.4
- [storage shelves](#), in §4.4
- [storage type](#), in §4.1
- [storage usage](#), in §6
- [types](#), in §4.1
- [usage](#), in §8

## § Terms defined by reference

- [ECMAScript] defines the following terms:
  - agent
  - agent cluster
- [HTML] defines the following terms:
  - Navigator
  - WorkerNavigator
  - browsing session (for environment settings object)
  - current settings object
  - environment settings object
  - in parallel
  - opaque origin
  - origin (for environment settings object)
  - queue a task
  - relevant settings object
- [INFRA] defines the following terms:
  - append
  - ascii string
  - clone
  - contain
  - exist
  - for each (for map)
  - implementation-defined
  - is empty
  - key
  - map
  - remove
  - set
  - string
  - user agent
  - value
- [PERMISSIONS] defines the following terms:
  - "granted"
  - "persistent-storage"
  - permission revocation algorithm
  - permission state
  - powerful feature
  - requesting permission to use
- [WEBIDL] defines the following terms:
  - Exposed
  - Promise
  - SameObject
  - SecureContext
  - TypeError
  - boolean
  - this
  - unsigned long long

## § **References**

### § **Normative References**

#### **[ECMAScript]**

[ECMAScript Language Specification](https://tc39.es/ecma262/). URL: <https://tc39.es/ecma262/>

#### **[HTML]**

Anne van Kesteren; et al. [HTML Standard](https://html.spec.whatwg.org/multipage/). Living Standard. URL:

<https://html.spec.whatwg.org/multipage/>

#### **[INFRA]**

Anne van Kesteren; Domenic Denicola. [Infra Standard](https://infra.spec.whatwg.org/). Living Standard. URL:

<https://infra.spec.whatwg.org/>

#### **[PERMISSIONS]**

Mounir Lamouri; Marcos Caceres; Jeffrey Yasskin. [Permissions](https://w3c.github.io/permissions/). URL:

<https://w3c.github.io/permissions/>

#### **[WEBIDL]**

Boris Zbarsky. [Web IDL](https://heycam.github.io/webidl/). URL: <https://heycam.github.io/webidl/>

## § IDL Index

```
[SecureContext]
interface mixin NavigatorStorage {
    [SameObject] readonly attribute StorageManager storage;
};
Navigator includes NavigatorStorage;
WorkerNavigator includes NavigatorStorage;

[SecureContext,
 Exposed=(Window,Worker)]
interface StorageManager {
    Promise<boolean> persisted();
    [Exposed=Window] Promise<boolean> persist();

    Promise<StorageEstimate> estimate();
};

dictionary StorageEstimate {
    unsigned long long usage;
    unsigned long long quota;
};
```